

NAVIGATION-BY-PREFERENCE: A NEW CONVERSATIONAL RECOMMENDER WITH PREFERENCE-BASED FEEDBACK

Group 5

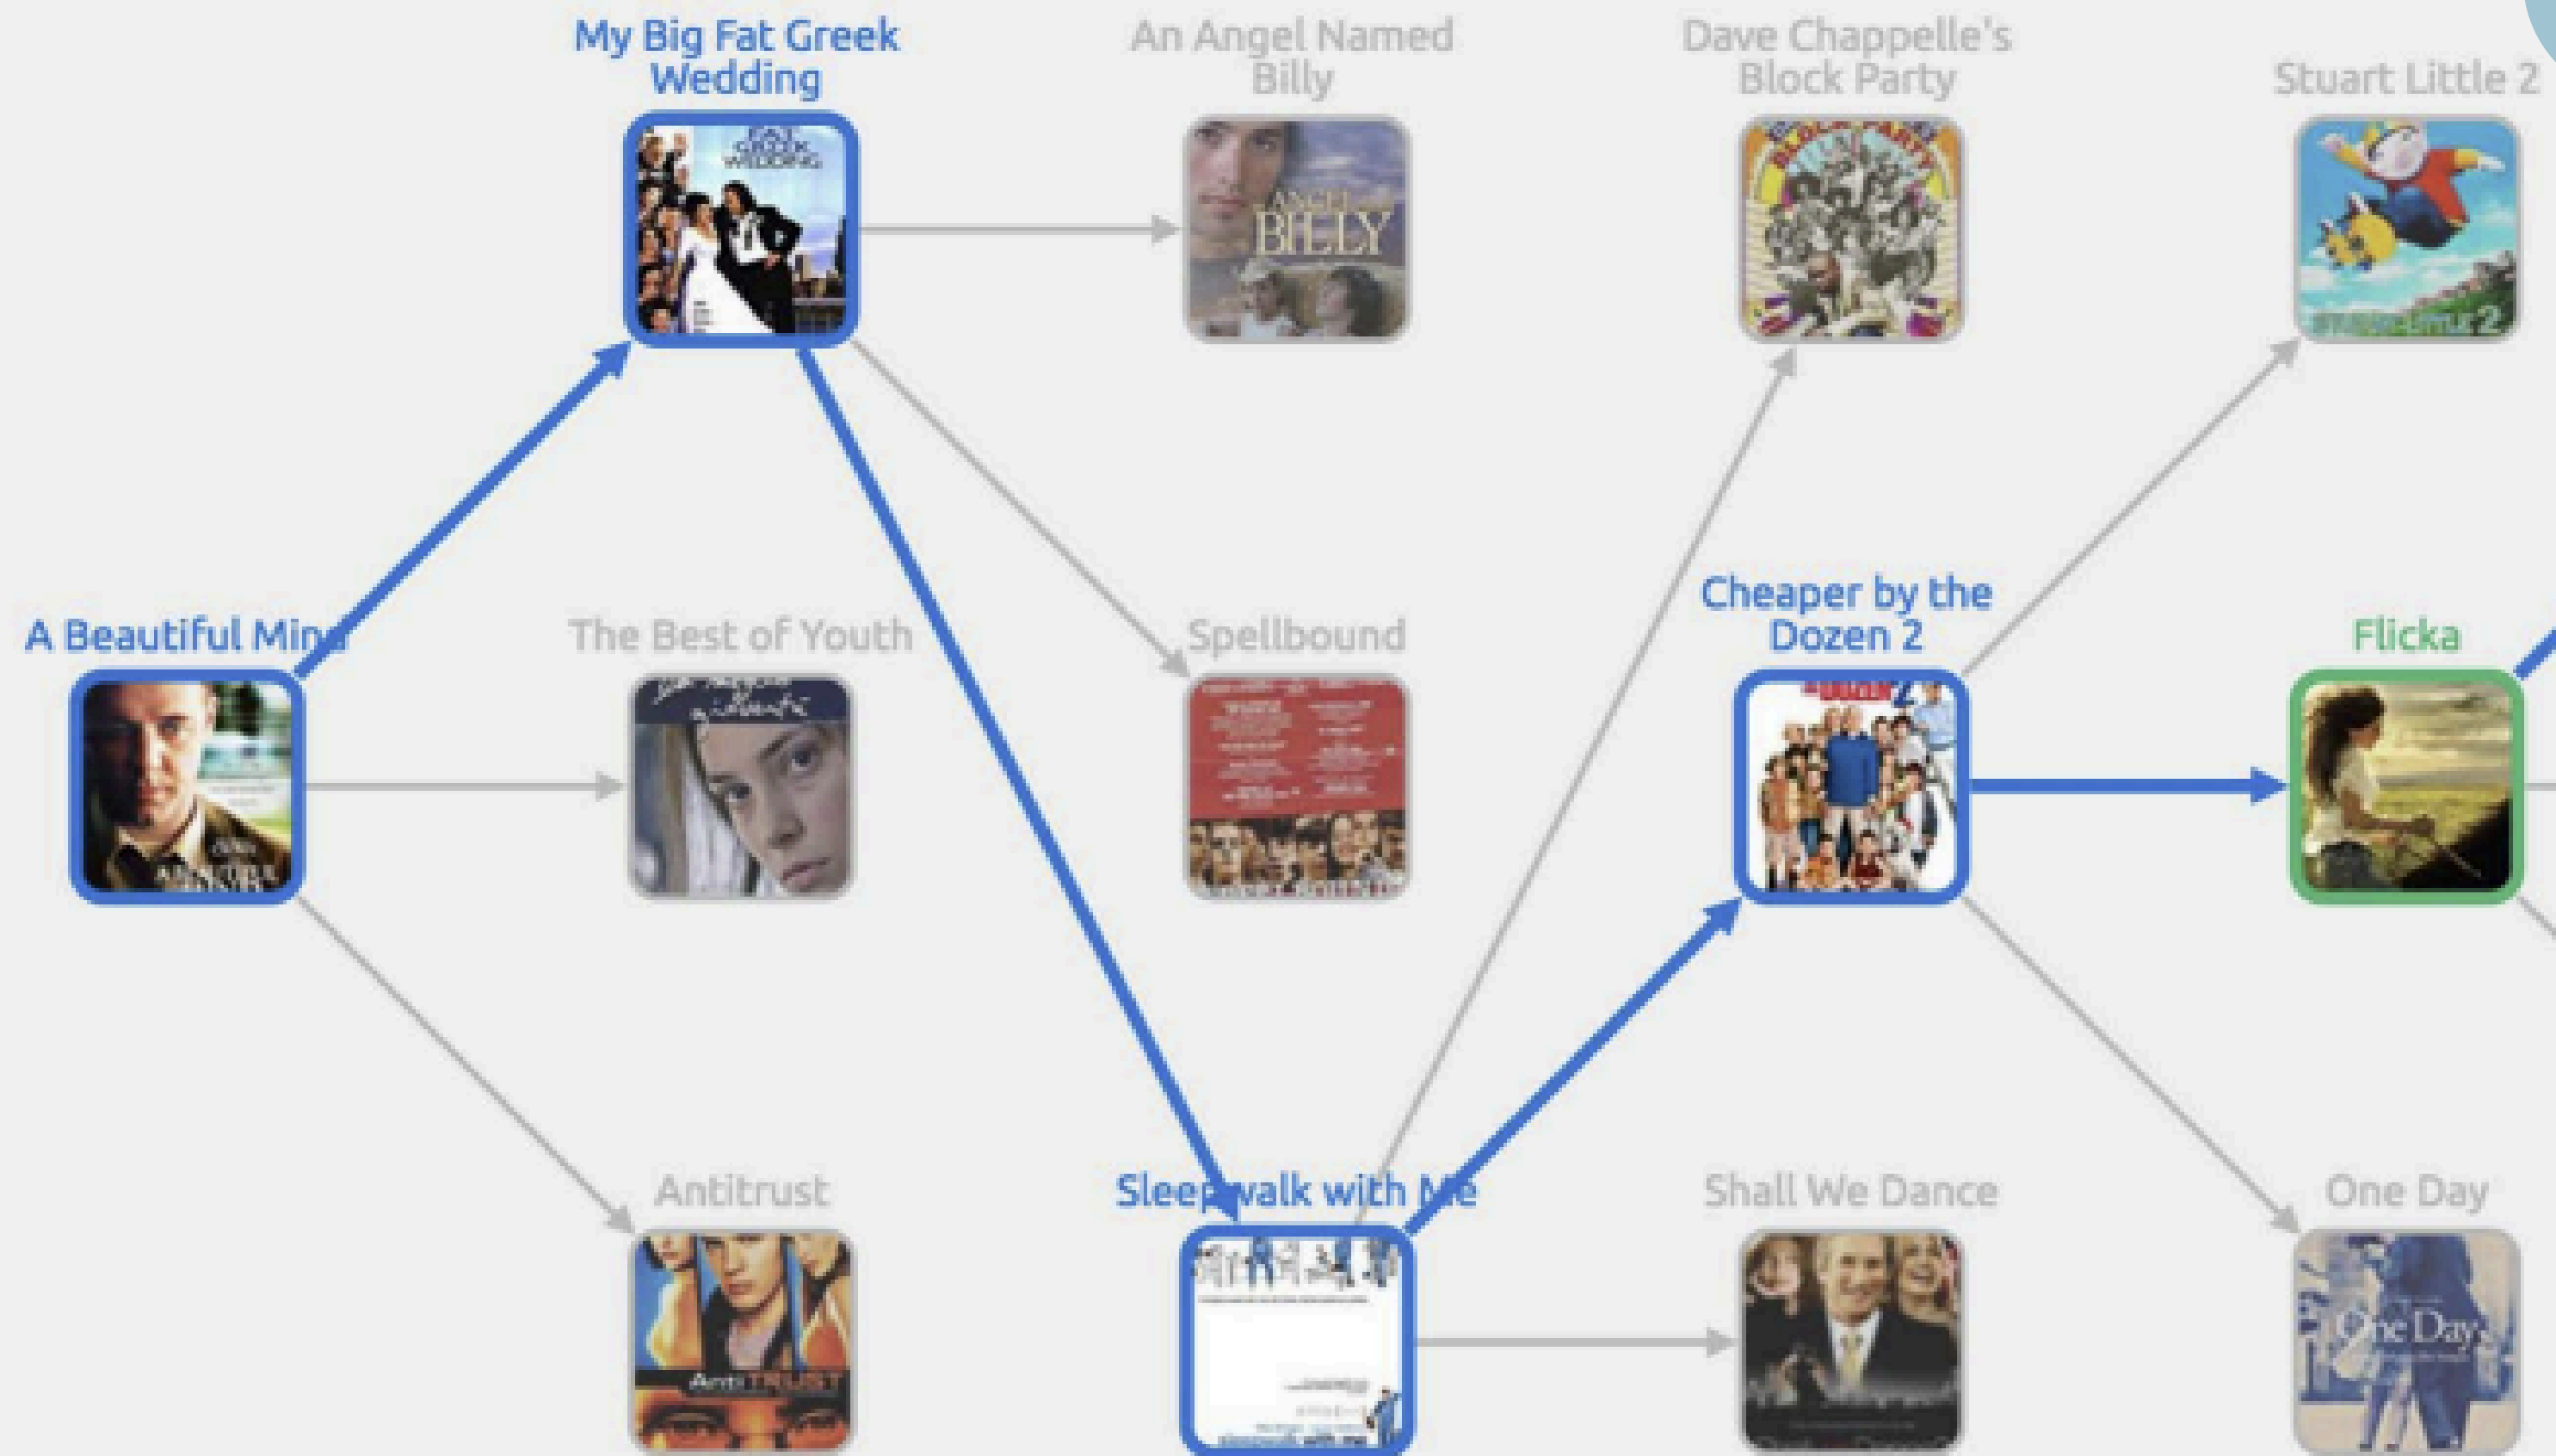
DAIICT | 2024

MOTIVATION

- **Traditional systems offer a single shot of recommendations.**
- **This can be limiting if users aren't satisfied with their top n recommendation**
- **(Eg = users with different short term and long term preferences)**
- **Hence, to overcome the drawback of single shot recommendations we use conversational recommenders.**

INTRODUCTION

- In a conversational recommendation system, a set of n recommendation is suggested
- Then the user is asked to choose one of the recommendations
- This item becomes query for next cycle
- May take several cycles



NOVELTIES

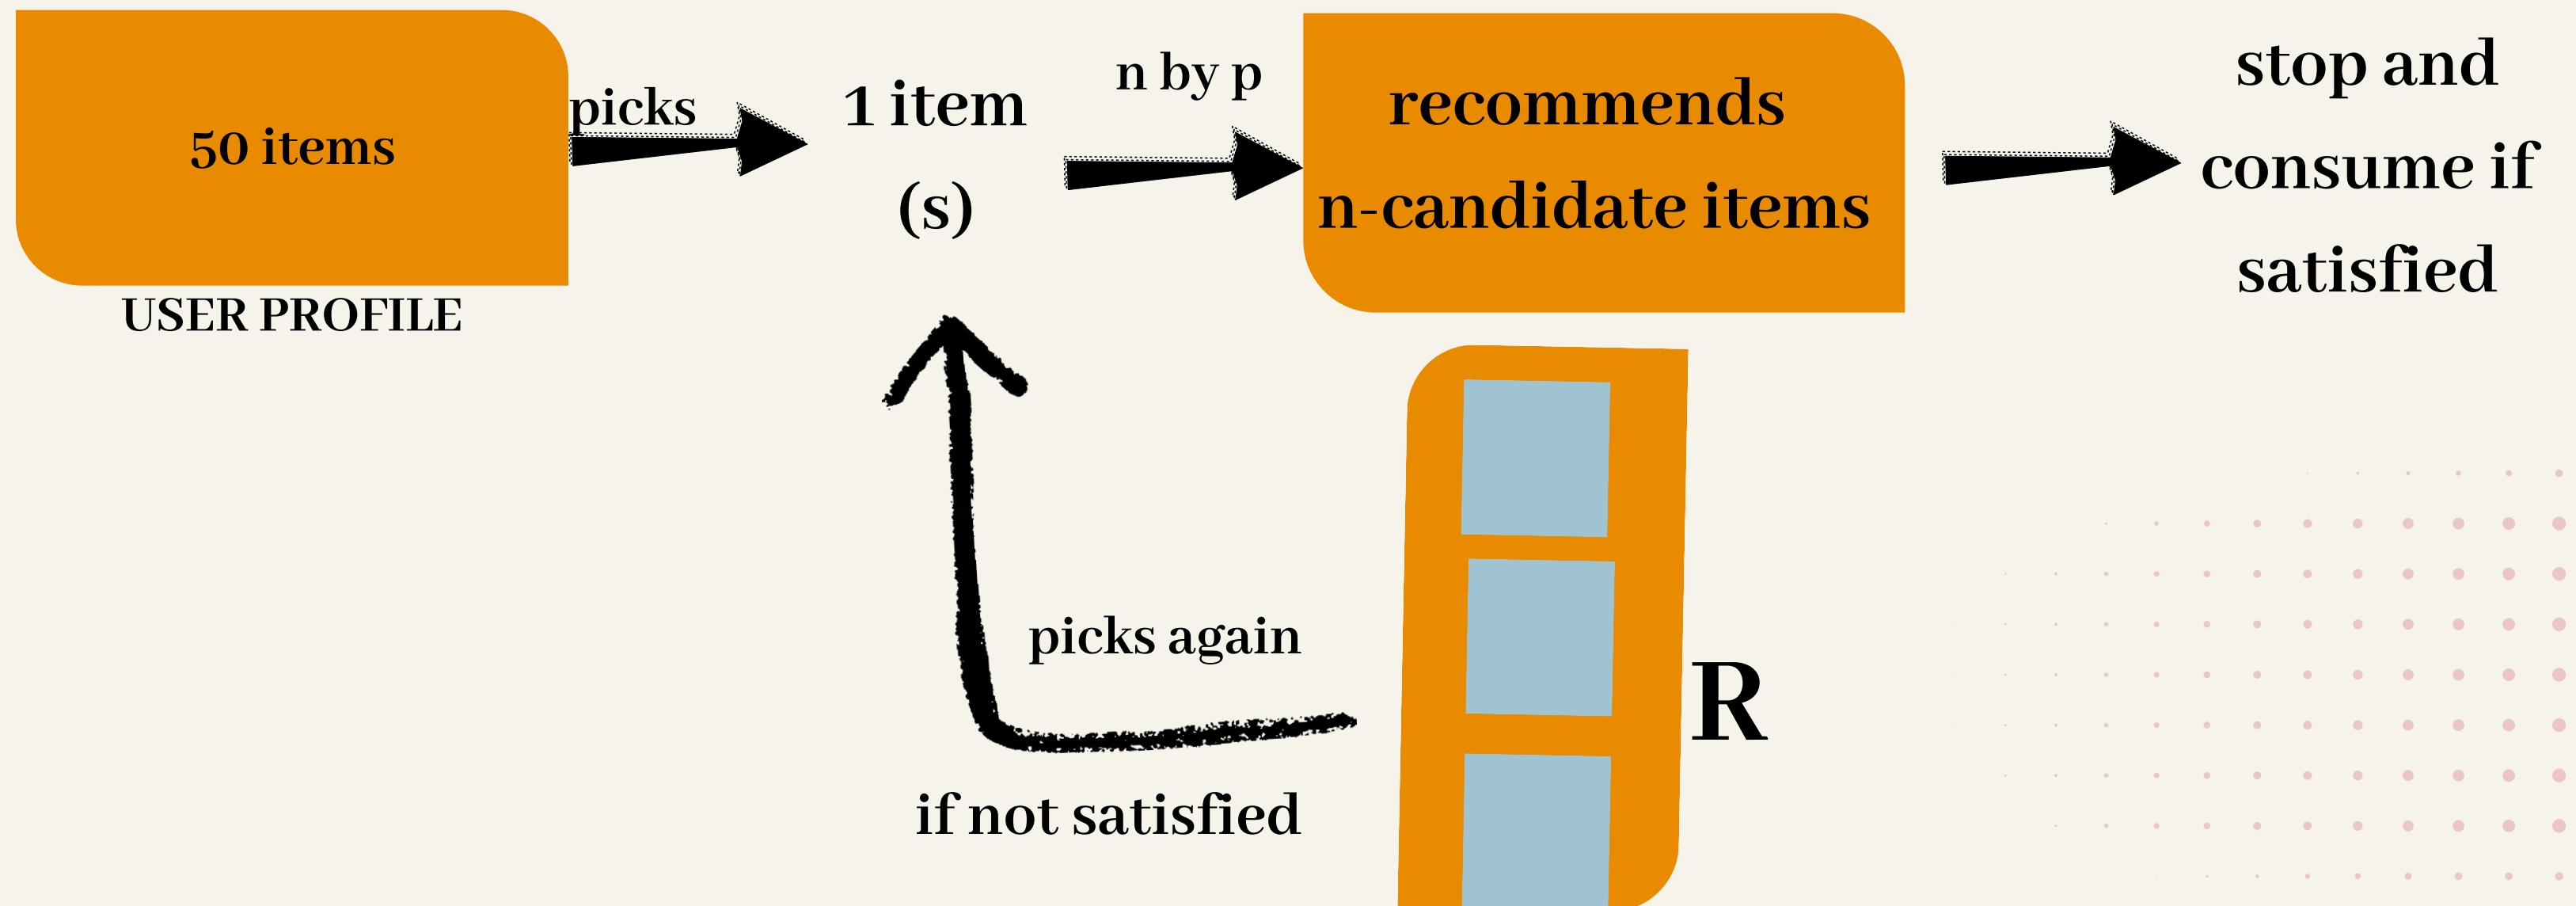
- **Works for unstructured feature values**
- **Takes into account the combination of short term as well as the long term goals instead of just short term goals**
- **Evaluation done combining both offline experiments and user trial**

NAVIGATION BY PREFERENCE

- It is a conversational recommendation system that used feedback based on user preferences (i.e short term /long term).
- The most important advantage of this system is its ability to work in domains that have unstructured item representations.
- It does not involve learning of models from the training sets of session data.

NAVIGATION BY PREFERENCE

High level working of n-by-p



NOTATIONS

P – user profile of active user
(set of items liked by user) $P \subseteq I$

I – set of all items (movies)

I – all items that might be recommended to user (candidate items).

$$I = \{i : i \in I \setminus P\}$$

f_i – set of features (tags) for item **i**

$$\text{sim}(i, j) = \frac{|f_i \cap f_j|}{|f_i \cup f_j|}$$

N_i – set of related items
(i.e candidate items that are neighbours of **i**)

$$N_i = \{j \in I, \quad j, i : \text{sim}(i, j) > \theta\}$$

WHICH ITEMS TO RECOMMEND IN NEXT CYCLE ?

- Candidates that are similar to s i.e N_s .
- But not every member of N_s should be recommended

WHY ?

- It is because we will not recommend an item more than once in whole dialog.
- Also the items not selected in previous round might be treated as negative feedback

- Since the user's most recent choice is $s \in R$
- To reflect negative feedback (i.e. $R \setminus \{s\}$), we have to discard some members of N_s .
- We refer to this subset of N_s as $S \subseteq N_s$. This is called the selection consistent candidate set.
- $L = \bigcup_{i \in P} N_i$ (i.e User's long term preferences)
(candidate items that are neighbours of each item i in user profile)

GOALS OF SET "R"

- **To reflect long term preferences of user**
(measure how much N_i overlaps with L)
- **To reflect short term preferences of user**
(measure how much N_i overlaps with S as a whole)
- **To ensure diversity**
(make sure that N_i covers parts of S that were not covered in previous recommendation cycles)

TYPES

n-by-i-p

In this, only feedback from the most recent cycle of recommendation affects the next cycle.

n-by-c-p

In this, feedback from the earlier cycles of recommendation also affects the next cycle.

N-BY-I-P

Algorithm 1 Navigation-by-Immediate-Preference (*n-by-i-p*)

Input: s : seed item, chosen by the user

L : candidate items that are neighbours of items in P

π : update policy

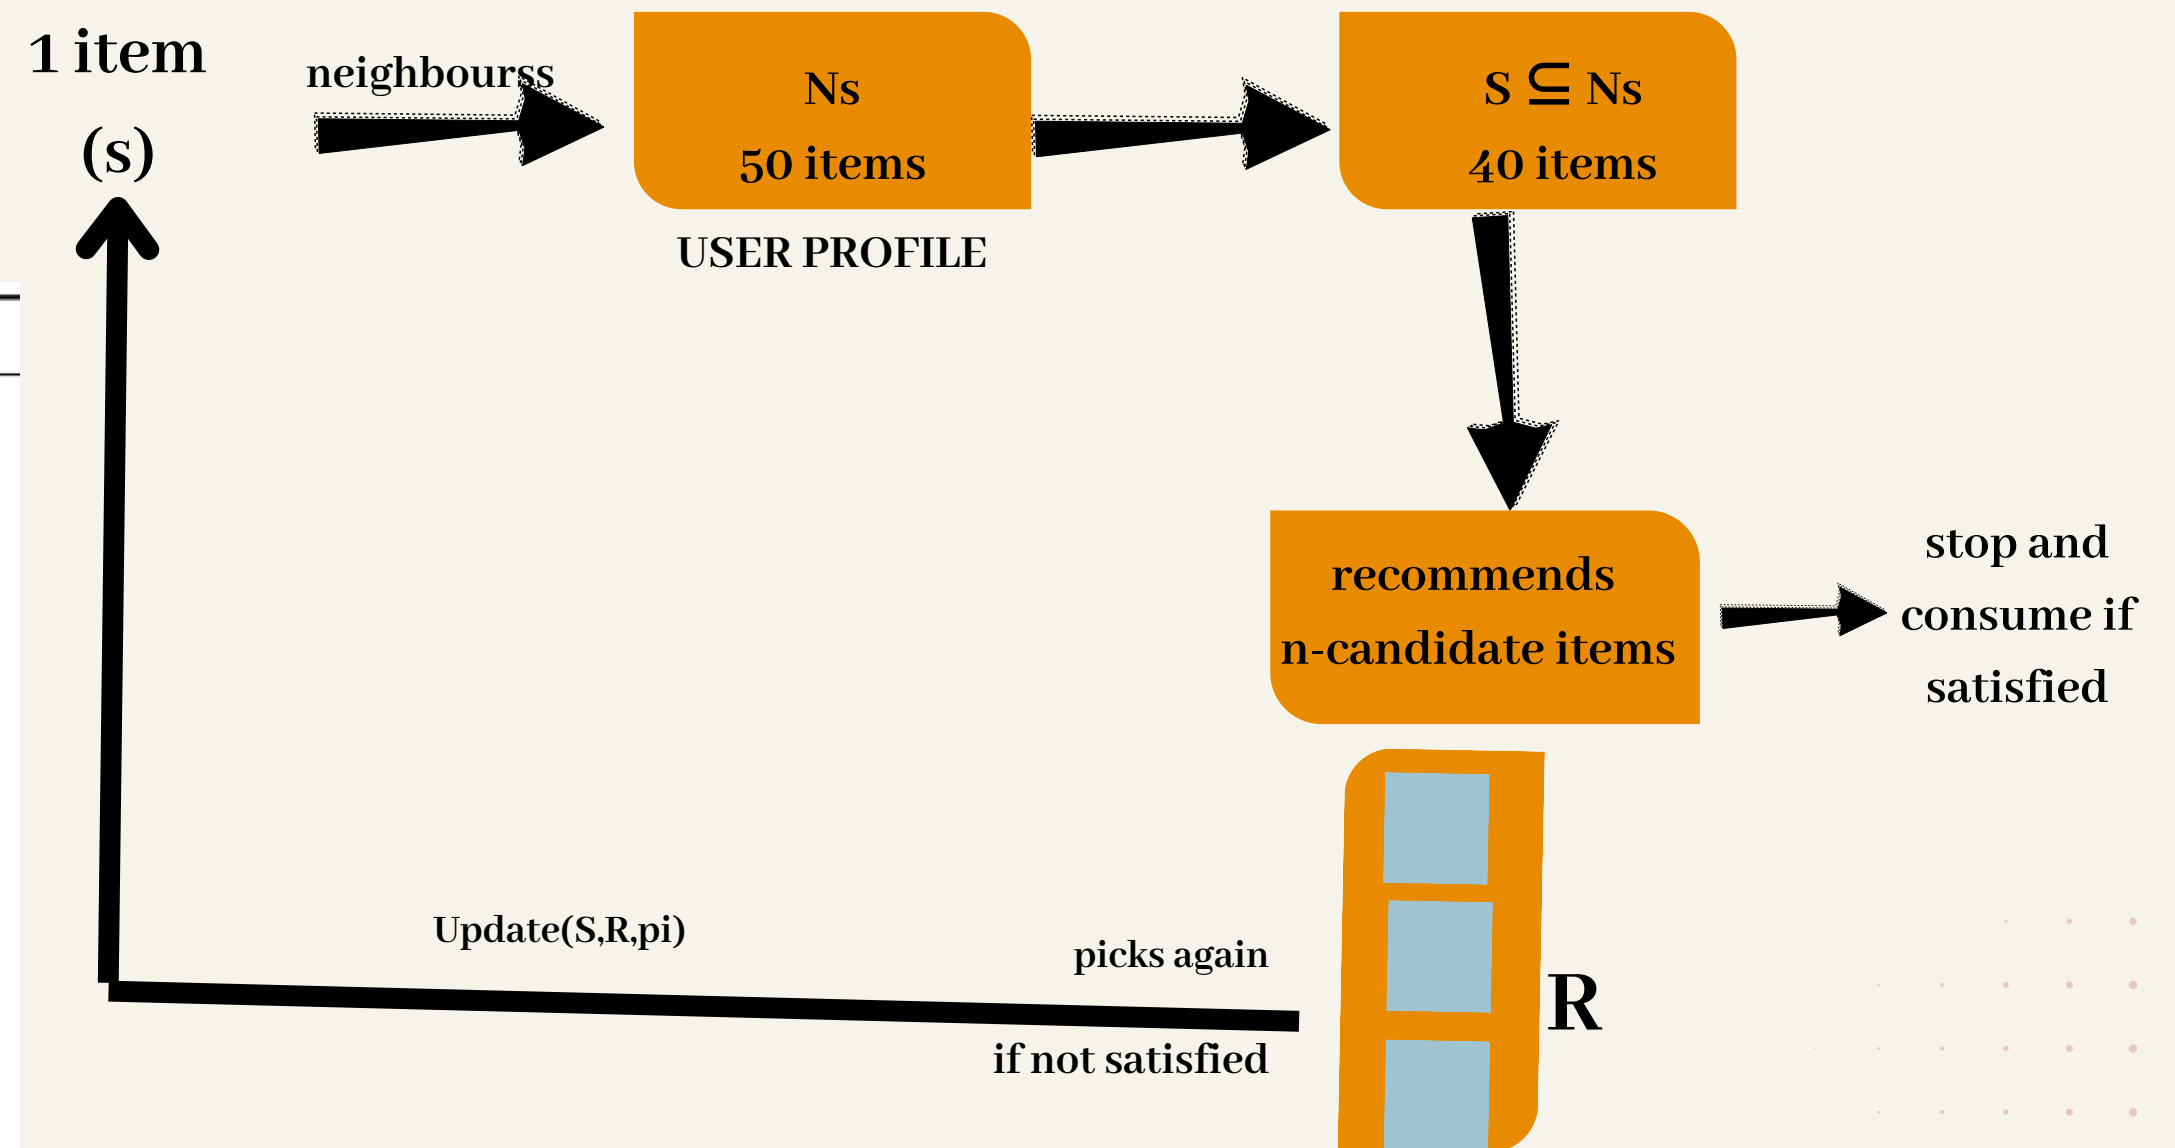
n : number of recommendations per cycle

Output: $i \in I$, a candidate item to consume

```

1:  $S \leftarrow N_s$ 
2:  $Tabu \leftarrow \emptyset$ 
3: while  $|S| > n$  do
4:    $R \leftarrow \text{RECOMMEND}(S, L, n)$ 
5:    $s, a \leftarrow$  user chooses  $s \in R$  and  $a \in \{STOP, CONTINUE\}$ 
6:   if  $a = STOP$  then
7:     return  $s$ 
8:    $S \leftarrow \text{UPDATE}(s, R \setminus \{s\}, \pi)$ 
9:    $Tabu \leftarrow Tabu \cup R$ 
10:   $S \leftarrow S \setminus Tabu$ 

```



RECOMMENDING

- Greedily selects n items from S with highest score

$$\text{score}(i, S, L, R) = (1 - \eta) \cdot \text{ovrlp}(i, S, R) + \eta \cdot \text{ovrlp}(i, L \setminus S, R)$$

prevents double counting of item in both S and L

$$\text{ovrlp}(i, X, R) = \frac{2 \cdot |(N_i \setminus \text{cov}(X, R)) \cap X|}{|N_i \setminus \text{cov}(X, R)| + |X \setminus \text{cov}(X, R)|}$$

to exclude items that are already covered by neighbours of R

No explicit reference to features

HOW TO UPDATE SET "S"

● Open

It simply means that we don't remove any members of set N_s .

i.e we don't take negative feedback into account.

● Strict

If any member of N_s is neighbour of any member of R' then we remove that member from N_s and the updated set N_s is assigned to S

i.e we don't recommend items from R' in next recommendation

● Relaxed

We remove member of N_s if it is neighbour of every (all) member of R'

HOW TO UPDATE SET "S"

15

● Mean

If similarity of member of N_s with s is less than MEAN of similarity of N_s with Rejected members then we remove those members from N_s


$$S \leftarrow N_s \setminus \{j \in N_s : \text{sim}(j,s) < \text{mean}(\text{Sims}(j))\}$$

● Max

If similarity of member of N_s with s is less than MAX of similarity of N_s with Rejected members then we remove those members from N_s

$$S \leftarrow N_s \setminus \{j \in N_s : \text{sim}(j,s) < \max(\text{Sims}(j))\}$$

PROBLEM WITH N-BY-I-P



Current set R may contain items which are not related to previously selected items by users.

Current set R may contain items which are related to previously rejected items by users.

N-BY-C-P

Algorithm 3 Navigation-by-Cumulative-Preference (*n-by-c-p*)

Input: s : seed item, chosen by the user

L : candidate items that are neighbours of items in P

ρ : re-weighting policy

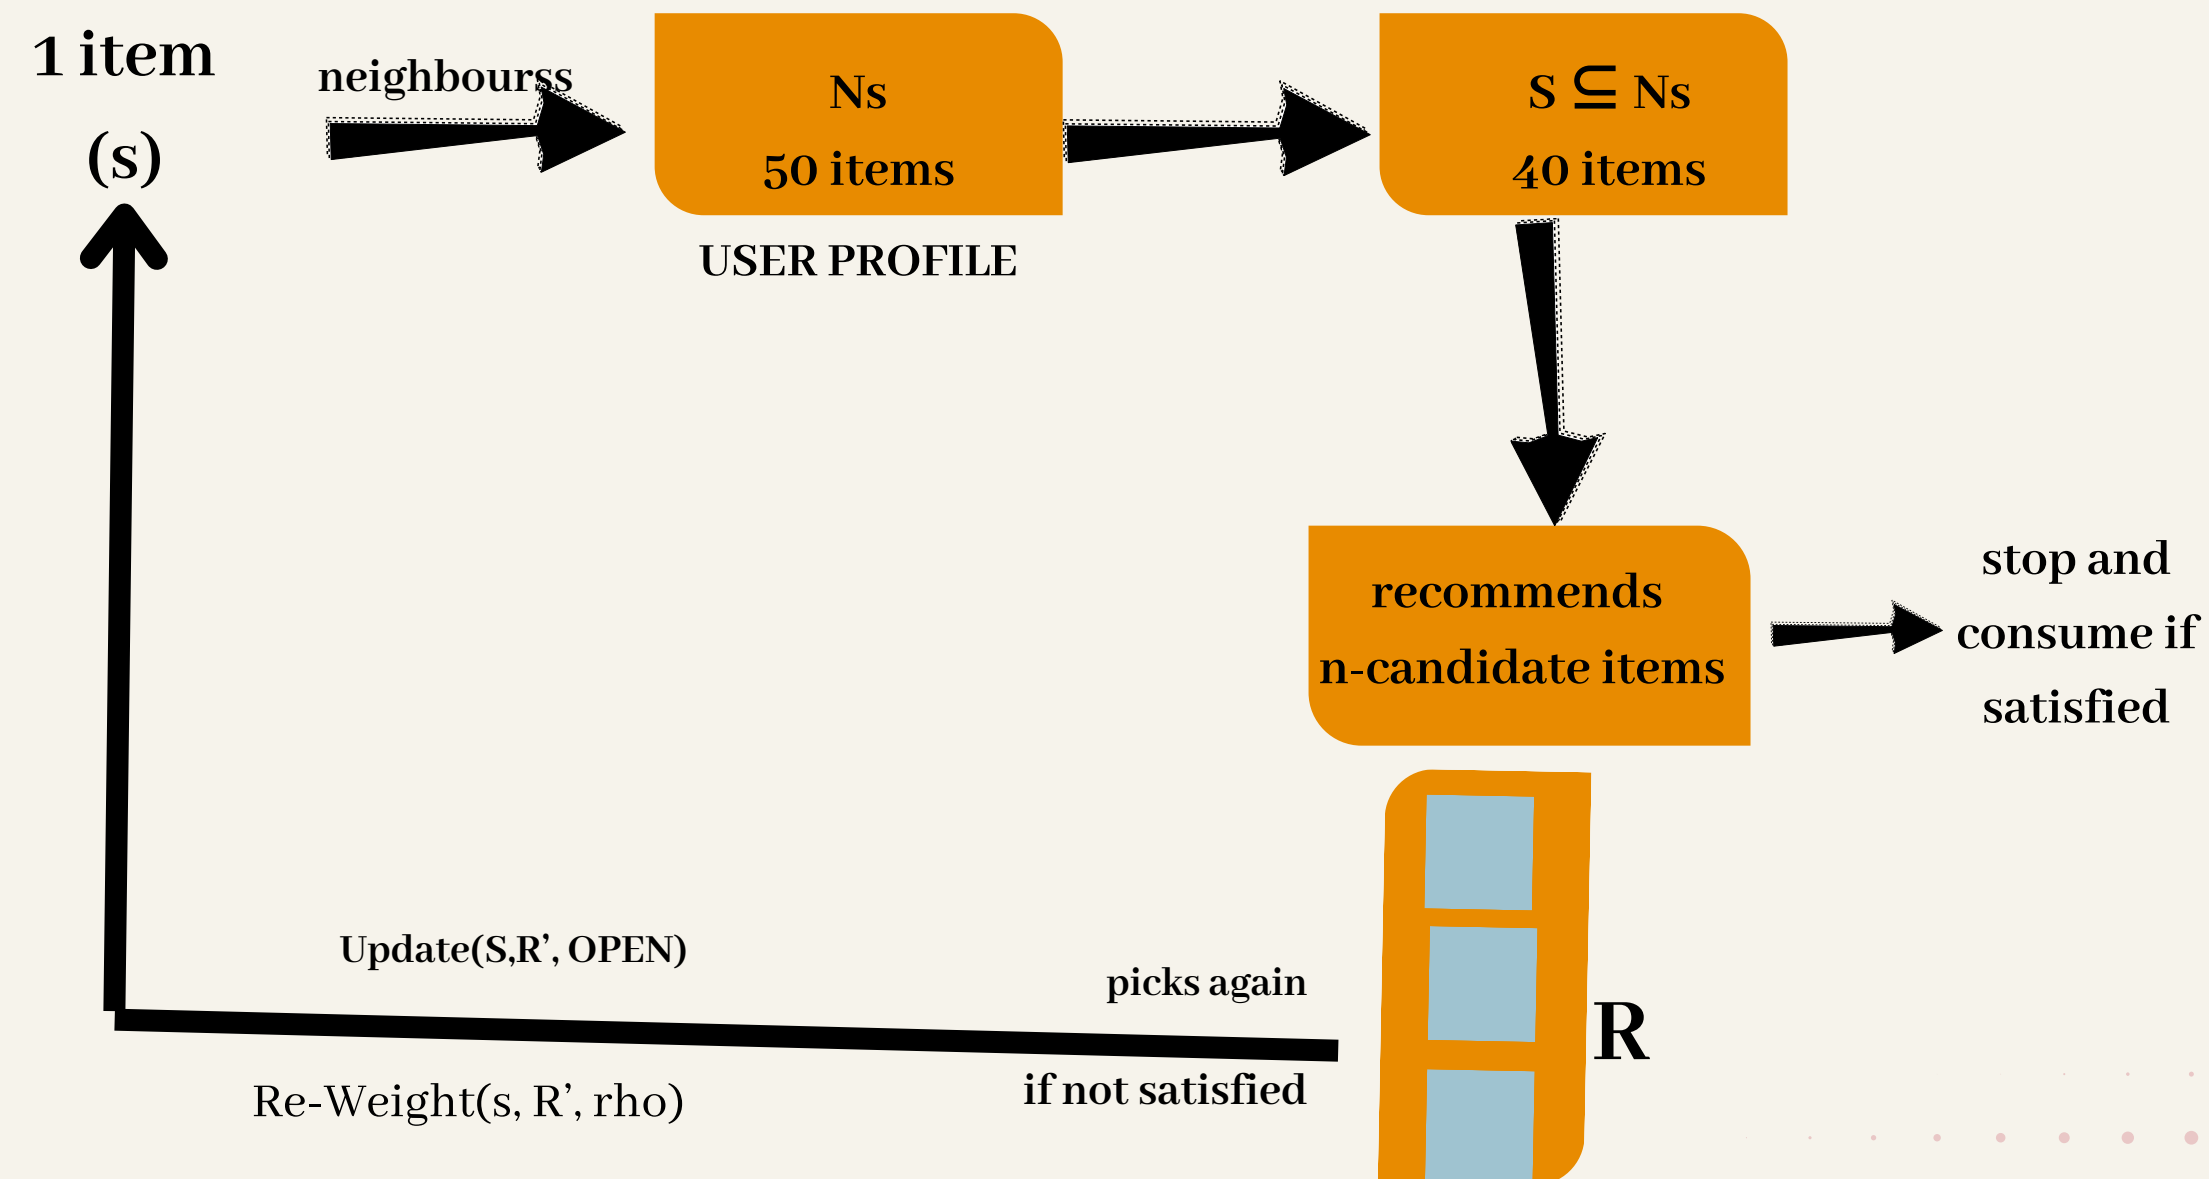
n : number of recommendations per cycle

Output: $i \in I$, a candidate item to consume

```

1:  $S \leftarrow N_s$ 
2:  $Tabu \leftarrow \emptyset$ 
3: REWEIGHT( $s, \emptyset, \rho$ )
4: while  $|S| > n$  do
5:    $R \leftarrow \text{RECOMMEND}(S, L, n)$ 
6:    $s, a \leftarrow$  user chooses  $s \in R$  and  $a \in \{STOP, CONTINUE\}$ 
7:   if  $a = STOP$  then
8:     return  $s$ 
9:    $S \leftarrow \text{UPDATE}(s, R \setminus \{s\}, \pi = \text{Open})$ 
10:  REWEIGHT( $s, R \setminus \{s\}, \rho$ )
11:   $Tabu \leftarrow Tabu \cup R$ 
12:   $S \leftarrow S \setminus Tabu$ 

```



RECOMMENDING

- Greedily selects n items from S with highest score

$$\text{wscore}(i, S, L, R) = (1 - \eta) \cdot \text{wovrlp}(i, S, R) + \eta \cdot \text{wovrlp}(i, L \setminus S, R)$$

prevents double counting of item in both S and L

$$\text{wovrlp}(i, X, R) = \frac{2 \cdot \sum_{j \in (N_i \setminus \text{cov}(X, R)) \cap X} w_j}{|N_i \setminus \text{cov}(X, R)| + |X \setminus \text{cov}(X, R)|}$$

insted of counting items directly, it counts the weights assigned to each item.

No explicit reference to features

RE-WEIGHTING

$w_i \leftarrow w_i + \Delta w_i \quad \forall i \in I$ (all candidate items)

$C_{ij} = \begin{cases} 1 & ; i \in N_s \\ 0 & ; \text{otherwise} \end{cases}$
check whether i and
neighbours of j are
related

Δw_i is Increased - when i is related to item just selected by user (C_{is})

Δw_i is Decreased - when i is related to item just rejected by user (C_{ij})

RE-WEIGHTING

$\rho = \text{Directional (Direc):}$

$$\Delta w_i = C_{is} - \sum_{j \in R'} C_{ij}$$

Only considers whether i is a neighbour of s or members of R' .

$\rho = \text{Similarity (Sim):}$

$$\Delta w_i = C_{is} \cdot \text{sim}(i, s) - \sum_{j \in R'} C_{ij} \cdot \text{sim}(i, j)$$

Considers similarities when i is a neighbour of s or members of R' .

$\rho = \text{Similarity-Mean (Smean):}$

$$\Delta w_i = C_{is} \cdot \text{sim}(i, s) - \text{mean}(\{C_{ij} \cdot \text{sim}(i, j) : j \in R'\})$$

Considers similarity when i is a neighbour of s , and the mean similarity when i is a neighbour of members of R' .

$\rho = \text{Similarity-Max (Smax):}$

$$\Delta w_i = C_{is} \cdot \text{sim}(i, s) - \max(\{C_{ij} \cdot \text{sim}(i, j) : j \in R'\})$$

Considers similarity when i is a neighbour of s , and the maximum similarity when i is a neighbour of members of R' .

$\rho = \text{Recency (Rcy):}$

$$\Delta w_i = C_{is} \cdot \text{sim}(i, s)^{1/d} - \sum_{j \in R'} C_{ij} \cdot \text{sim}(i, j)^{1/d}$$

As per *Sim* above, but with updates counting more for later recommendations.

$\rho = \text{Recency-Mean (Rmean):}$

$$\Delta w_i = C_{is} \cdot \text{sim}(i, s)^{1/d} - \text{mean}(\{C_{ij} \cdot \text{sim}(i, j)^{1/d} : j \in R'\})$$

As per *Smean* above, but with updates counting more for later recommendations.

$\rho = \text{Recency-Max (Rmax):}$

$$\Delta w_i = C_{is} \cdot \text{sim}(i, s)^{1/d} - \max(\{C_{ij} \cdot \text{sim}(i, j)^{1/d} : j \in R'\})$$

As per *Smax* above, but with updates counting more for later recommendations.

NOTION OF 'JUMP'

- It is possible that no item in current set R suits to user but the one from previous cycle is more suitable.
- User can select either to consume that previous item or create a new set of recommendations based on that item.

OFFLINE EXPERIMENTS

The goal for the Offline Experiments conducted are:

- "n-by-i-p" versus "n-by-c-p": These represent different ways of considering user feedback, either focusing on the most recent feedback or all feedback across dialog cycles.
- Different update and re-weighting policies for handling negative feedback.
- The influence of a parameter called η , which controls the balance between short-term and long-term user preferences.
- These experiments were used to decide which configurations to use in user trial.

EXPERIMENT SETTINGS

- The hetrec dataset included over 2000 users, 6000 movies, 80000 keywords and half a million ratings.
- 500 randomly chosen users are taken and profiles were created for users based on movies they liked (rated 4 or 5).
- Simulated dialogs were created between users and different configurations of the n-by-p system (maximum cycles for recommendations =15).

HOW TO CAPTURE SHORT TERM PREFERENCES

24

- Randomly choose a target item t from N_s



- Recommend n -most similar items to target t



- Stop the simulated dialog when $t \in R$
(i.e target item is one of the recommendation.)



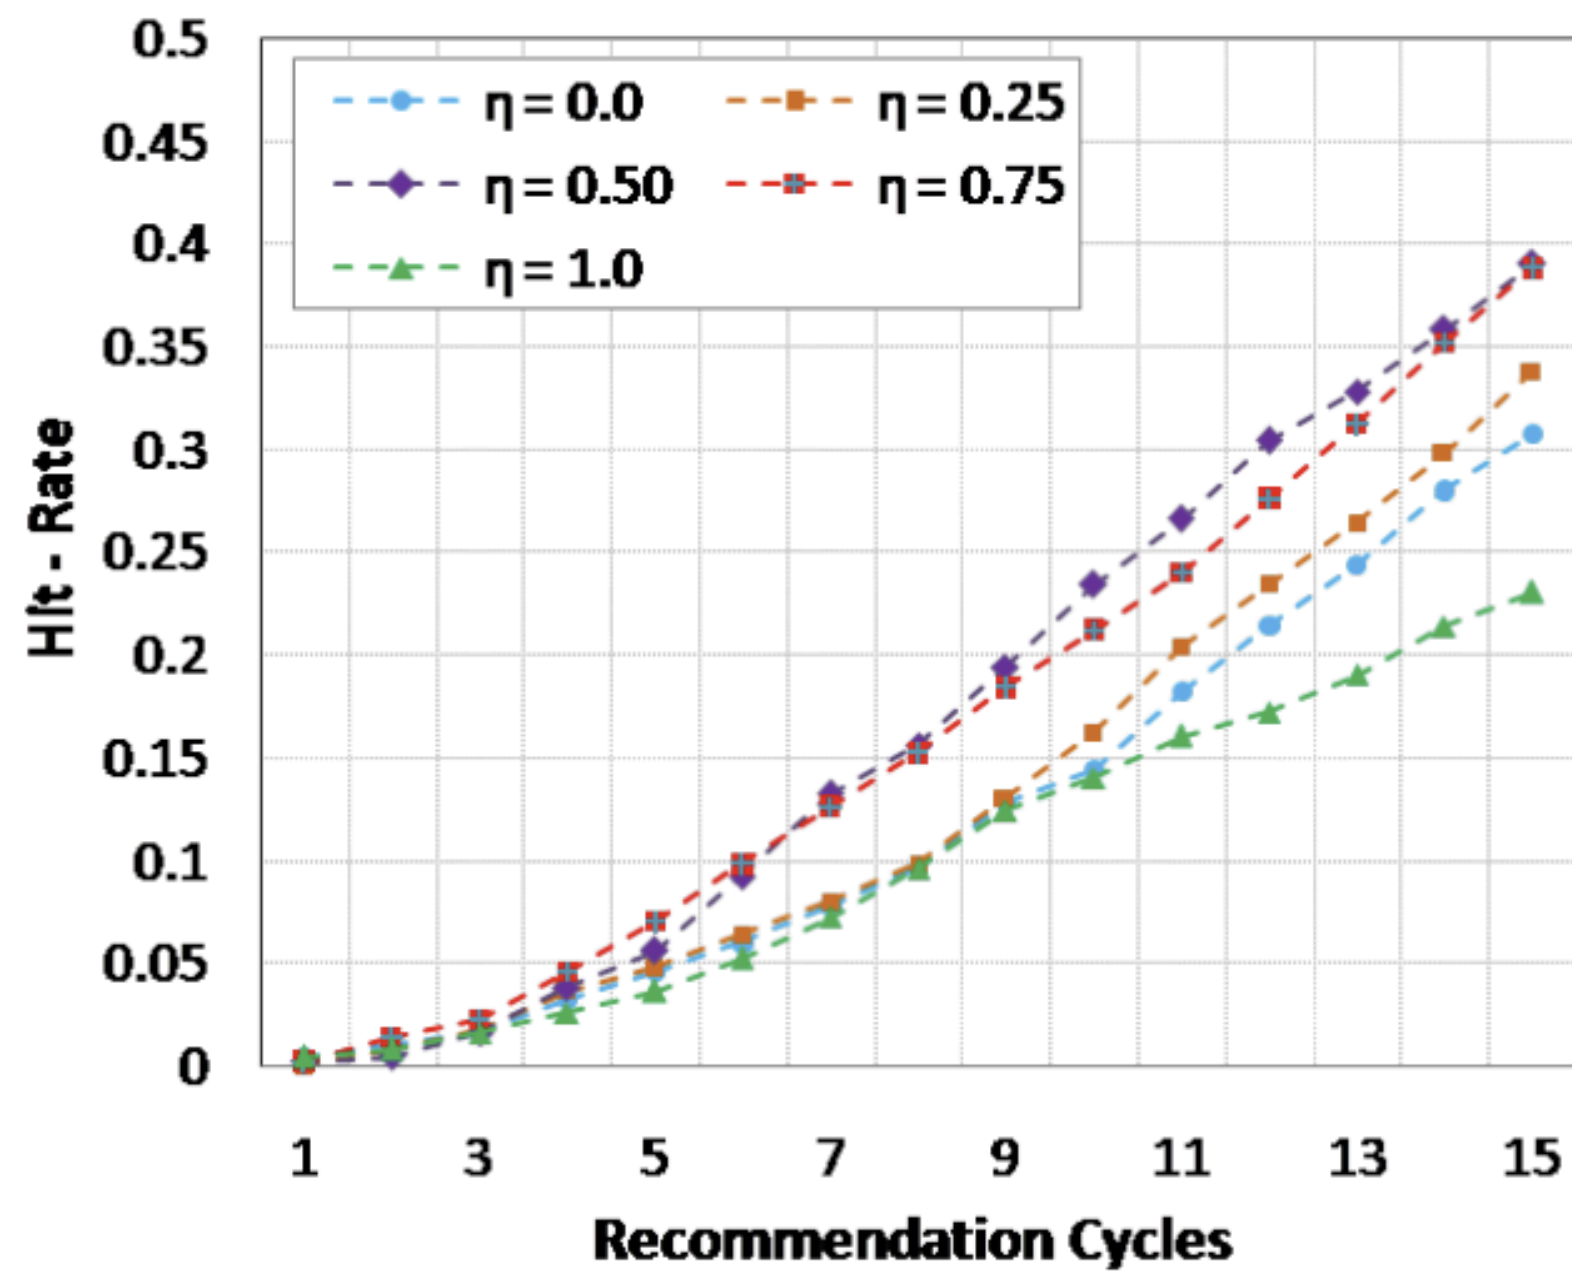
Measure Hit rate for each cycle

EXPERIMENT METRICS


- **Hit-rate:** The proportion of simulated users who found their target item within 15 cycles.
- **Jaccard similarity:** A measure of similarity between the item selected by the simulated user and their target.
- **Diversity and surprise:** Metrics indicating the variety and unexpectedness of recommendations (Div_cont and S_cont).

EXPERIMENT RESULTS

η	<i>n-by-i-p</i>					<i>n-by-c-p</i>						
	<i>Strict</i>	<i>Relaxed</i>	<i>Open</i>	<i>Mean</i>	<i>Max</i>	<i>Direc</i>	<i>Sim</i>	<i>Smean</i>	<i>Smax</i>	<i>Rcy</i>	<i>Rmean</i>	<i>Rmax</i>
0.00	0.102	0.042	<u>0.204</u>	0.146	0.098	0.112	0.070	0.308	0.282	0.028	0.204	0.174
0.25	0.058	0.034	0.152	0.072	0.044	0.112	0.040	0.338	0.270	0.016	0.196	0.158
0.50	0.032	0.046	0.036	0.038	0.050	0.084	0.040	<u>0.390</u>	0.250	0.018	0.222	0.154
0.75	0.016	0.038	0.014	0.014	0.042	0.072	0.042	0.388	0.266	0.018	0.162	0.134
1.00	0.012	0.030	0.010	0.024	0.042	0.070	0.040	0.230	0.220	0.020	0.112	0.122



(a) Hit-rate in each cycle



(b) Mean Jaccard similarity in each cycle between the user's selected item and their target item

Figure 2: Results per cycle for n -by- c - p with $\rho = S_{mean}$ and different values of η for *random* targets.

● Purpose

To evaluate the effect of incorporating long-term preferences alongside short-term preferences in recommendation accuracy and user satisfaction.

● Systems Compared

- **Smean@0.5:** Utilizes a mix of short-term and long-term preferences ($\eta = 0.5$) as the best-performing setup from previous offline experiments.
- **Smean@0.0 Baseline:** Similar to Smean@0.5 but exclusively focuses on short-term preferences ($\eta = 0.0$), serving as the control group.

TRIAL DATA

- A total of 102 people participated in trial.
- Trial used a subset of the dataset from previous offline experiments.
- Featured 1851 movies released between 2000 and 2011, to maximize participant familiarity.
- Random user assignment to either $S_{mean@0.5}$ or $S_{mean@0.0}$.
- Goal: Assess how the integration of long-term preferences affects user interactions and system performance.

USER TRIAL PROTOCOL

30

- Participants begin by selecting 10 movies they like, establishing their long-term preferences.



- Users select a movie from their profile they think is best to watch with a selected family member with differing tastes.



- Users select from a list of eight family members for whom they think movie preferences differ.
- They then choose a movie they believe both would enjoy.
- Process may be repeated up to three times to refine the ephemeral goal.

Please select a person you know but whose movie preferences are different from yours.

Mother

Father

Brother

Sister

Aunt

Uncle


Nephew

Niece

You are going to choose a movie to watch with your Father

Choose a movie from your profile that you would enjoy watching together. (If none of them are ideal, choose the least worst one.)

Shrek
2001 | 90 min



A Beautiful Mind
2001 | 135 min



How much do you think
you and your Father will
enjoy watching this movie
together?

- ☐ Not at all
- ☐ Barely at all
- ☐ Fair
- ☐ Somewhat
- ☐ A lot

The Pursuit of Happ...
2006 | 117 min



Street Fight
2005 | 83 min



Figure 3: Parts of screenshots showing selection of a companion and a seed movie.

TRIAL

- **Eight-cycle dialog where the system presents three movie recommendations in each cycle.**
- **After eight cycles, users are presented with a tree of 24 movies.**
- **They select the movie believed to best suit the viewing with their companion.**
- **Maintains consistency across trials, allowing for fair comparison despite possible discovery of an 'ideal' movie early in the session.**

USER EXPERIMENT

- **Evaluation Metrics:**

1. **Familiarity:** Whether users had seen the movie before.
2. **Relevance:** Users' judgment on potential enjoyment of the movie with a companion.
3. **Serendipity:** Surprise element of the recommendation.
4. **Effectiveness:** Helpfulness of the recommendations.
5. **Satisfaction:** Overall enjoyment of using the system.

- **Statistical Analysis:**

One-sided Z-test and t-test used to determine significance of differences.

USER TRIAL RESULTS

Responses to survey questions in the user trial.

User's Response	<i>Smean@0.5</i>				<i>Smean@0.0</i>			
	<i>Relevance</i>	<i>Serendipity</i>	<i>Effectiveness</i>	<i>Satisfaction</i>	<i>Relevance</i>	<i>Serendipity</i>	<i>Effectiveness</i>	<i>Satisfaction</i>
<i>Not at all</i>	0	2	0	1	3	5	5	2
<i>Barely at all</i>	1	4	5	5	7	4	4	5
<i>Fair</i>	11	12	14	12	16	21	17	14
<i>Somewhat</i>	18	19	21	20	12	12	16	18
<i>A lot</i>	21	14	11	13	13	9	9	12

IMPROVEMENT FORMULA

Formula:

$$\text{improvement} = \frac{\sum_{v=1}^5 v \cdot \omega_v - \sum_{v=1}^5 v \cdot \alpha_v}{\sum_{v=1}^5 5 \cdot \omega_v - \sum_{v=1}^5 v \cdot \alpha_v}$$

- Measures how much the system improved the satisfaction from the initial movie to the recommended final movie.
- A higher value indicates more significant improvement relative to the best possible outcome.

USER EFFORT

Measure of effort	<i>Smean@0.5</i>	<i>Smean@0.0</i>
<i>Nodes displayed</i>	26.85	29.56
<i>Node mouse-overs</i>	19.28	22.84
<i>Edge mouse-overs</i>	10.03	11.92
<i>Cycles</i>	5.36	4.12
<i>Time taken (secs.)</i>	251.93	300.58

All values are averaged over participants who liked their final selected movie Somewhat or A lot.

CONCLUSIONS

- **Preference based feedback does not use features explicitly.**
- **n-by-p is highly configurable (i.e 60 configurations)**
- **The system is interpretable and easy to understand the relationship between pairs of consecutive items in preference chain.**
- **User efforts are minimal.**

The background features three vertical stripes on the left: a wide pink stripe, a medium blue stripe, and a narrow beige stripe. The right side of the slide is white, adorned with a grid of small pink dots that fades out towards the right edge.

THANK YOU

Presented By :

Shubham Gupta (202318052)

Kavisha Madani (202318007)

Shashwat Parikh (202318062)

Mayan Bhut (202318043)